

# RNA-Seq Pre-processing

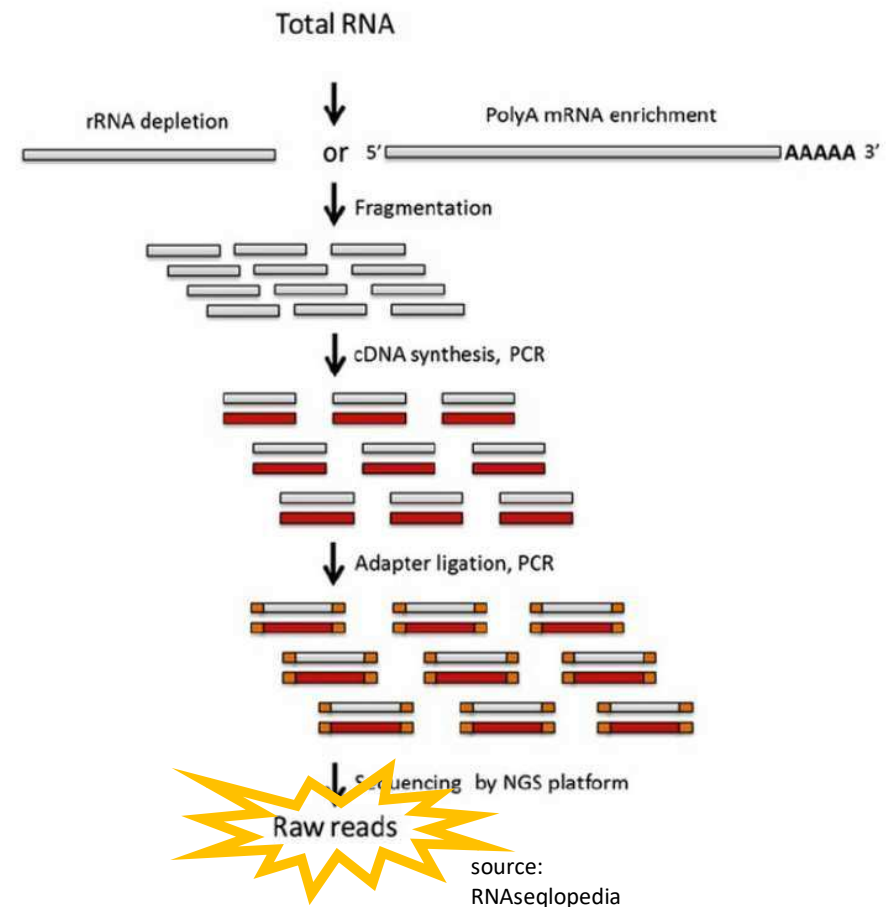
Lauren Vanderlinden

BIOS 6660

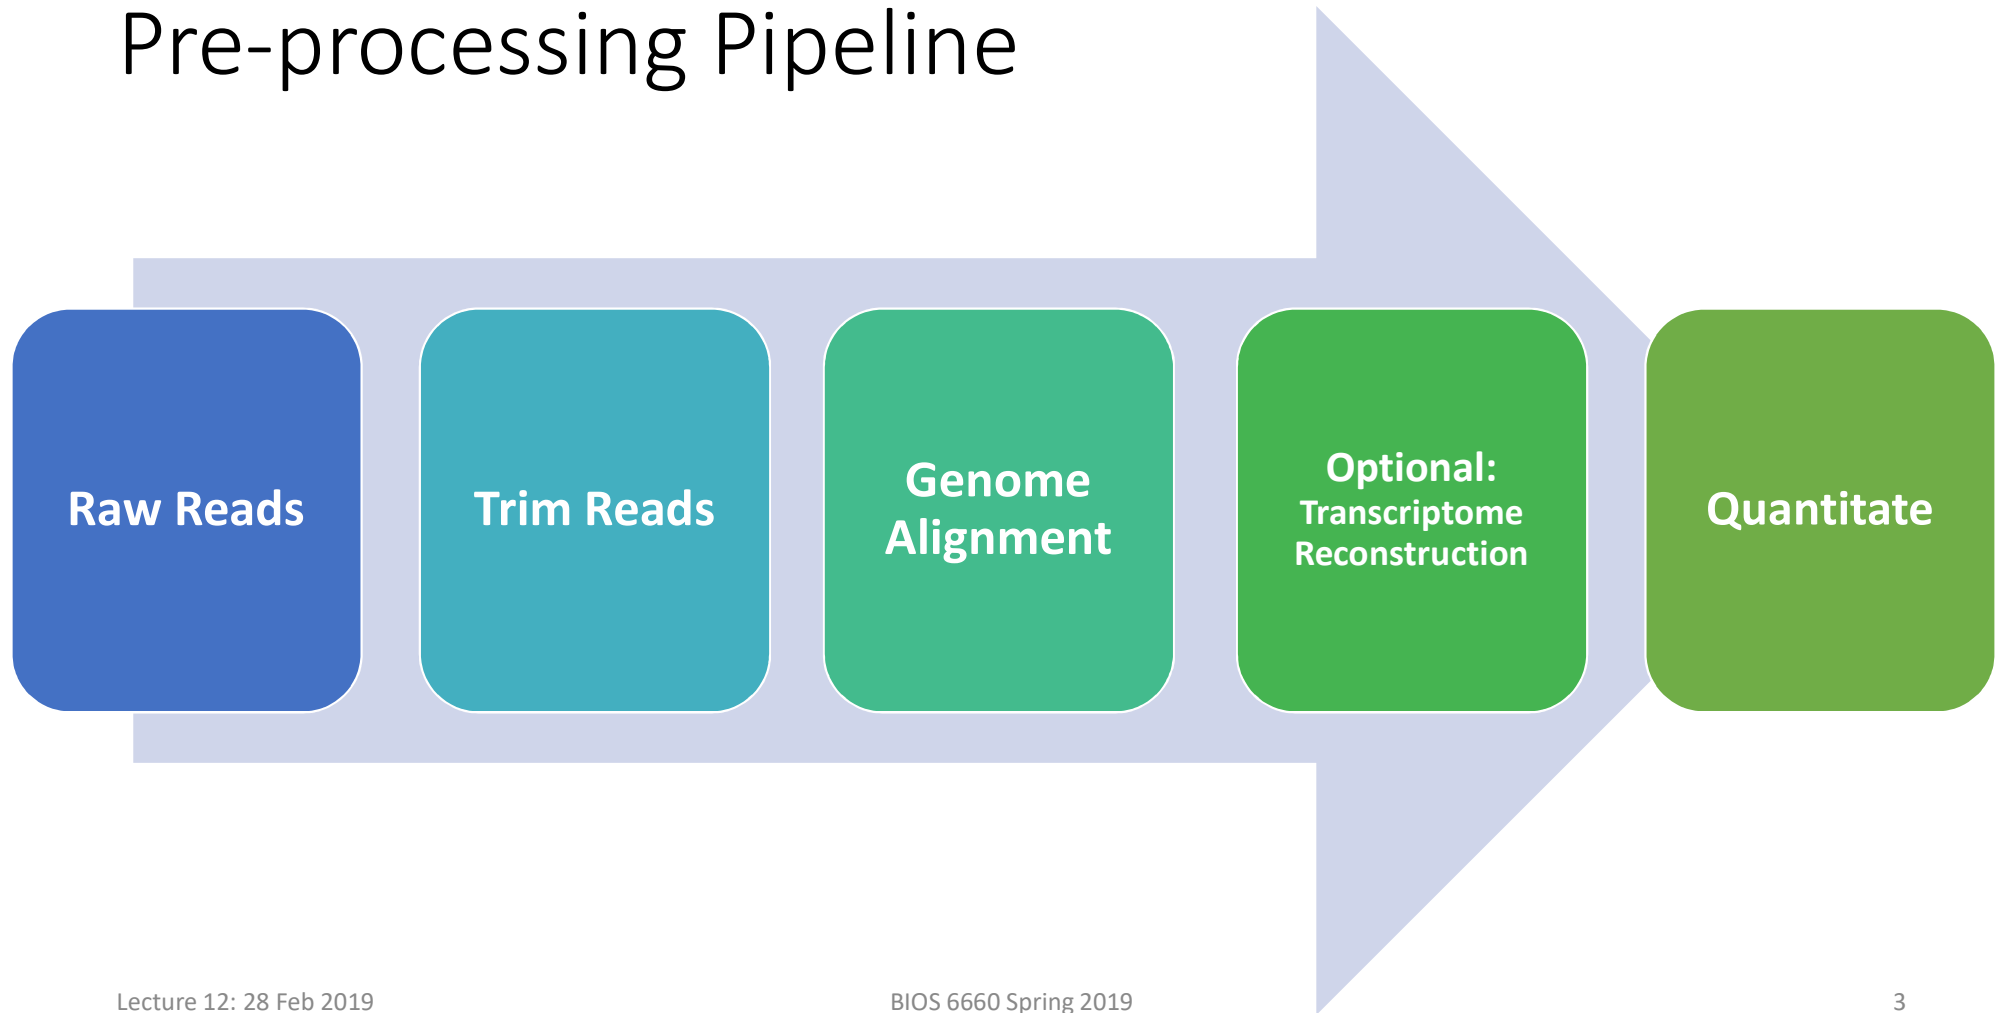
Spring 2019

# Overview from Last Time

- Talked about RNA-Sequencing
  - Importance of RNA expression
  - Technology used to get sequences
- 
- Now you have raw reads!



# Pre-processing Pipeline



# Folder Structure

- For a typical coverage where the main goal is differential expression (25-30 million reads/sample) raw files range around 2-3.8 GB
- Find it useful to have these files in separate folders for each part of the process:
  - rawReads
  - trimmedReads
  - alignedReads
  - quantitation
  - data (matrix form finally)
- Suggestion only

/data/home/vanderll/Britt/TagSeq.20181206/

Name

- ↑ ..
- rawReads
- bigwig
- trimmedReads
- quantitation
- alignedReads
- code
- reports
- data

A horizontal arrow pointing right, divided into five colored segments: blue (labeled 'Raw Reads'), light blue, teal, green, and dark green.

ID	@NS500358:142:HLKWMBGX3:1:11101:14573:1049 1:N:0:CTATAC
Sequence	GATAANACATGTAAGCCCGATTGTAGGCAGATTTTGGCCTTTTTGTTCGCAACAATGTTGATTACTTAATGACGCACTTAT GTTGGCCCGTGCAATCCATGAATCGGAAGATCGGAAGAGCACACGTCTGAACTCCAGTCACCTATACATCT
Quality	+ AAAAA#EEEEEAEEEEEE<EEEEAEAEAA EEEEEEEEEEEEEEEEFAEEEEEEEEEEEEEEEEEEEEFAAE<E<<EEE/EEE6AEEEEEEEEAAE<//// @NS500358:142:HLKWMBGX3:1:11101:25309:1049 1:N:0:CTATAC

- Lecture 12: 28 Feb 2019 BIOS 6660 Spring 2019 5

# FASTQ Illumina ID Format

```
@HWUSI-EAS100R:6:73:941:1973#0/1
```

<b>HWUSI-EAS100R</b>	the unique instrument name
<b>6</b>	flowcell lane
<b>73</b>	tile number within the flowcell lane
<b>941</b>	'x'-coordinate of the cluster within the tile
<b>1973</b>	'y'-coordinate of the cluster within the tile
<b>#0</b>	index number for a multiplexed sample (0 for no indexing)
<b>/1</b>	the member of a pair, /1 or /2 ( <i>paired-end or mate-pair reads only</i> )

Source: wikipedia

# Phred Quality (Q) Score

- A score assigned to each individual base called
- Developed to help in the automation of DNA sequencing in the Human Genome Project
- Phred score is logarithmically related to the base-calling error probabilities

Phred quality scores are logarithmically linked to error probabilities

Phred Quality Score	Probability of incorrect base call	Base call accuracy
10	1 in 10	90%
20	1 in 100	99%
30	1 in 1000	99.9%
40	1 in 10,000	99.99%
50	1 in 100,000	99.999%
60	1 in 1,000,000	99.9999%

Illumina standards say Phred of 30 or higher considered “high quality”

Source: wikipedia

# FASTQ and Phred Quality Score

- FASTQ score represented as ASCII-33 characters

```
vanderll@sysgen:~>gzip -cd /data/home/sabal/Britt/RNA-Seq.2017-11-06/rawReads/A1
_S01_S1_R1_001.fastq.gz | head
@Ns500358:142:HLKWMGX3:1:11101:14573:1049 1:N:0:CTATAC
GATAANACATGTAAGCCAGTTTAGGCAGATTTGCCTTTTGTTCGCAACAAATGTTGATTACTTAATGACGCACTTAT
GTTGGCCGTGCAATCCATGAATCGGAAGATCGGAAGAGCACACGTCTGAACTCCAGTCACCTATACATCT
+
AAAAA#EEEEEEEEEE<EEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE
EEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE<E<<EEE/EEE6AAAAEEEEAE<////
@Ns500358:142:HLKWMGX3:1:11101:25309:1049 1:N:0:CTATAC
```

Quality

ASCII\_BASE=33 Illumina, Ion Torrent, PacBio and Sanger

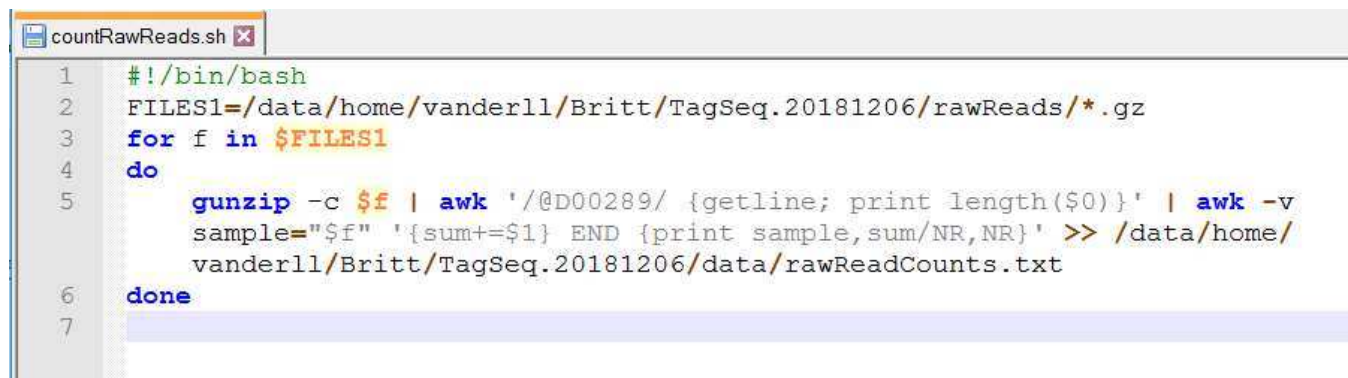
Q	P_error	ASCII	Q	P_error	ASCII	Q	P_error	ASCII	Q	P_error	ASCII
0	1.00000	33 !	11	0.07943	44 ,	22	0.00631	55 7	33	0.00050	66 B
1	0.79433	34 "	12	0.06310	45 -	23	0.00501	56 8	34	0.00040	67 C
2	0.63096	35 #	13	0.05012	46 .	24	0.00398	57 9	35	0.00032	68 D
3	0.50119	36 \$	14	0.03981	47 /	25	0.00316	58 :	36	0.00025	69 E
4	0.39811	37 %	15	0.03162	48 0	26	0.00251	59 ;	37	0.00020	70 F
5	0.31623	38 &	16	0.02512	49 1	27	0.00200	60 <	38	0.00016	71 G
6	0.25119	39 '	17	0.01995	50 2	28	0.00158	61 =	39	0.00013	72 H
7	0.19953	40 (	18	0.01585	51 3	29	0.00126	62 >	40	0.00010	73 I
8	0.15849	41 )	19	0.01259	52 4	30	0.00100	63 ?	41	0.00008	74 J
9	0.12589	42 *	20	0.01000	53 5	31	0.00079	64 @	42	0.00006	75 K
10	0.10000	43 +	21	0.00794	54 6	32	0.00063	65 A			

Source: USEARCH



# How many reads do I have?


- One of the first questions before anything else is how many reads do my samples have?
- Make a bash script to count the number of reads you have




```
countRawReads.sh
1  #!/bin/bash
2  FILES1=/data/home/vanderll/Britt/TagSeq.20181206/rawReads/*.gz
3  for f in $FILES1
4  do
5      gunzip -c $f | awk '/@D00289/ {getline; print length($0)}' | awk -v
      sample="$f" '{sum+= $1} END {print sample, sum/NR, NR}' >> /data/home/
      vanderll/Britt/TagSeq.20181206/data/rawReadCounts.txt
6  done
7
```

- Notice I'm counting the number of lines which start with @D00289, not the total lines in the file
- Writes a text file with your fastq name and number of reads

# Trim Reads



- Adapter Trimming
  - Trimming off the adapter sequences of reads
  - ABSOLUTE MUST for small RNA
  - Improves *de novo* assemblies
- Quality Trimming
  - May increase mapping rates
  - May also lose information
- Software Programs
  - BBDuk
  - Cutadapt
  - Trim Galore!
  - PRINSEQ
  - Trimmomatic
  - Sickle/Sythe
  - FASTX Toolkit



Source: SciLifeLab

# Cutadapt code

- Cutadapt is on yampa

Extra options I've used:

- q quality score cut-off before adapter removal
- m minimum length if read to keep after trimming
- u remove a specified length of bases from beginning

```
Usage:
  cutadapt -a ADAPTER [options] [-o output.fastq] input.fastq

For paired-end reads:
  cutadapt -a ADAPT1 -A ADAPT2 [options] -o out1.fastq -p out2.fastq in1.fastq
  in2.fastq
```


- Example bash script

```
trimmReads.v3.sh
1  #!/bin/bash
2  FILES1=/data/home/sabal/Britt/RNA-Seq.2017-11-06/rawReads/*_R1_001*.fastq.gz
3  for f in $FILES1
4  do
5      f2=${f//R1/R2}
6      f_trimmed=${f//.fastq.gz/_trimmed.fastq.gz}
7      f_trimmed=${f_trimmed//rawReads/trimmedReads}
8      f2_trimmed=${f2//.fastq.gz/_trimmed.fastq.gz}
9      f2_trimmed=${f2_trimmed//rawReads/trimmedReads}
10     cutadapt -u 15 -U 15 -q 20 -m 20 -a AGATCGGAAGAGCACACGTCTGAACTCCAGTCAC -A AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGTAGATCTCGGTGGTCGCCGTATCATT -o $f_trimmed -p $f2_trimmed $f $f2
11 done
```

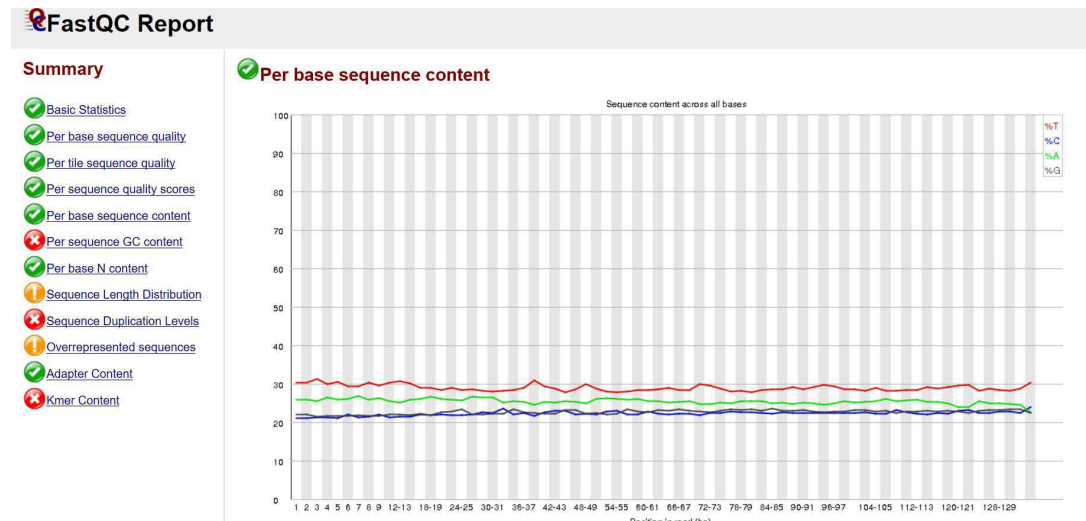

# FastQC

- FastQC installed on yampa
- Program which evaluates quality of reads
- Use this to check the trimming went well
- Want to stay in the green zone
- Really want green checks on:
  - Per sequence quality scores
  - Per base sequence content
- Look at example report

```
fastqc /path/sample.fastq.gz
```



# FastQC Pre and Post Quality Trimming




```
cutadapt -q 20 -m 20 -a AGATCGGAAGAGCACACGTCTGAACTCCAGTCAC -A AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGTAGATCTCGGTGGTCGCCGTA
TCATT -o $f_trimmed -p $f2_trimmed $f $f2
```

Removing an extra 15 bases from start

```
cutadapt -u 15 -U 15 -q 20 -m 20 -a AGATCGGAAGAGCACACGTCTGAACTCCAGTCAC -A AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGTAGATCTCG
GTGGTCGCCGTCATCATT -o $f_trimmed -p $f2_trimmed $f $f2
```

# Align Reads to Genome



- Align reads to a reference genome
  - Humans: Genome Reference Consortium human genome build 38 (GRCh38, aka hg38)
    - hg19 is still commonly used

Species	Most recent reference genome
Fruit fly	dm6
Mouse	mm10
Rat	rn6
Human	hg38

# Alignment Tools

## 1. Unspliced Aligners (No Gaps Allowed)

- BWA
- Bowtie/Bowtie2
- NovoAlign, SeqMap, SEAL


## 2. Spliced Aligners (Allows Splice-Junctions)

- Erange
  - SpliceSeq
  - BBMap
  - Hisat/Hisat2
  - STAR
  - TopHat
- 
- Annotated Guided Aligners
- De Novo Splice Aligners

# Hisat2 Reference Genome

- Hisat2 is on yampa
- For any alignment tool, you need reference files

```
hisat2-build reference.fa reference.hisat
```



<https://uswest.ensembl.org/info/data/ftp/index.html>

Show 10 entries			
★	Species	DNA (FASTA)	cDNA (FASTA)
Y	<a href="#">Human</a> <i>Homo sapiens</i>	<a href="#">FASTA</a>	<a href="#">FASTA</a>
Y	<a href="#">Mouse</a> <i>Mus musculus</i>	<a href="#">FASTA</a>	<a href="#">FASTA</a>
Y	<a href="#">Zebrafish</a> <i>Danio rerio</i>	<a href="#">FASTA</a>	<a href="#">FASTA</a>
	<a href="#">Agassiz's desert tortoise</a>	<a href="#">FASTA</a>	<a href="#">FASTA</a>

/BIOS6660/Homework6/indexFiles/				
Name	Size	Changed	R...	O...
dm6.hisat.1.ht2	50...	12/4/2017 9:35:4...	r...	v...
dm6.hisat.2.ht2	34...	12/4/2017 9:35:4...	r...	v...
dm6.hisat.3.ht2	22 ...	12/4/2017 9:33:2...	r...	v...
dm6.hisat.4.ht2	34...	12/4/2017 9:33:2...	r...	v...
dm6.hisat.5.ht2	75...	12/4/2017 9:36:0...	r...	v...
dm6.hisat.6.ht2	35...	12/4/2017 9:36:0...	r...	v...
dm6.hisat.7.ht2	1 KB	12/4/2017 9:33:2...	r...	v...
dm6.hisat.8.ht2	1 KB	12/4/2017 9:33:2...	r...	v...



# Hisat2 Alignment

```
hisat2 -x /BIOS6660/Homework6/indexFiles -1 sample_1.fq.gz -2 sample_2.fq.gz | samtools view -bS - > alignedSample.bam
```


- Options you might want:
  - --un writing an output file for those that don't align
  - -N # mismatches you will allow (default 0)
- For paired-end reads, you will now go from 2 files/sample down to 1 file/sample
- BAM output file is a compressed binary version of a SAM file
- SAM = Sequence Alignment/Map format

Message it prints after running:

```
20000 reads; of these:
  20000 (100.00%) were unpaired; of these:
    1247 (6.24%) aligned 0 times
    18739 (93.69%) aligned exactly 1 time
    14 (0.07%) aligned >1 times
93.77% overall alignment rate
```

# Batch Alignment

- Python script called: alignBatch.genome.hisat.py
- Written by Spencer Mahaffey (thank you!)
- This gives you 4 files:
  - Aligned.bam
  - Splice\_junct.bed
  - Summary.txt
  - Unmapped.bam
- Look at py script



Spencer Mahaffey  
smahaffey

```
python /data/home/vanderl1/teaching/BIOS6660_spring2019/programs/alignBatch.genome.hisat.py --input-suffix .fastq.gz --index-dir /data/home/vanderl1/annotation/dm6.hisat2.reference -P /data/home/vanderl1/teaching/BIOS6660_spring2019/data/hw6/subsetFASTQpairs
```

# Alignment Error Example

- I got this following error after alignment one time
- Saying that there is a read that has more bases than scores assigned to it
- It did output a bam file up to this point
- What the issue was that the core uploaded the fastq files to the web and at some point their connection was interrupted so the fastq file was only partially uploaded

```
[samopen] SAM header is present: 1870 sequences.  
Error: Read D00289:29:CCT5JANXX:5:2316:4056:36315  
1:N:0:CAGGCG+CTTGTA has more read characters than  
quality values.  
terminate called after throwing an instance of 'int'  
(ERR): hisat2-align died with signal 6 (ABRT) (core  
dumped)
```

# BAM Files: Alignment Output

```
samtools view sample.sorted.bam | head
```

```
@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r001 99 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTG *
r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
r003 0 ref 9 30 5S6M * 0 0 GCCTAAGCTAA * SA:Z:ref,29,-,6H5M,17,0;
r004 0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC *
r003 2064 ref 29 17 6H5M * 0 0 TAGGC * SA:Z:ref,9,+,5S6M,30,1;
r001 147 ref 37 30 9M = 7 -39 CAGCGGCAT * NM:i:1
```

Source: samtools manual

Col	Field	Type	Regexp/Range	Brief description
1	QNAME	String	[!-?A-Z]{1,254}	Query template NAME
2	FLAG	Int	[0, 2 <sup>16</sup> - 1]	bitwise FLAG
3	RNAME	String	\* [:rname:^\*] [:rname:]*	Reference sequence NAME <sup>9</sup>
4	POS	Int	[0, 2 <sup>31</sup> - 1]	1-based leftmost mapping POSition
5	MAPQ	Int	[0, 2 <sup>8</sup> - 1]	MAPping Quality
6	CIGAR	String	\* ([0-9]+[MIDNSHPX=])+	CIGAR string
7	RNEXT	String	\* = [:rname:^\*] [:rname:]*	Reference name of the mate/next read
8	PNEXT	Int	[0, 2 <sup>31</sup> - 1]	Position of the mate/next read
9	TLEN	Int	[-2 <sup>31</sup> + 1, 2 <sup>31</sup> - 1]	observed Template LENgth
10	SEQ	String	\*  [A-Za-z-.]+	segment SEQUENCE
11	QUAL	String	[!-~]+	ASCII of Phred-scaled base QUALity+33

- Optional header '@'
- Alignment info
- Example: first read starts at chr4:  
1,076,030

[illegible]

# SAM Flags

<https://www.samformat.info/sam-format-flag>

## SAM Format

SAM Flag

SAM Flag (single)

Base Quality

Header

Alignment Tags

Ambiguity Codes

About

Please enter a SAM flag value here: e.g. 99



Paired: 0

Reset

#	Decimal	Description of read
1	1	Read paired
2	2	Read mapped in proper pair
3	4	Read unmapped
4	8	Mate unmapped
5	16	Read reverse strand
6	32	Mate reverse strand
7	64	First in pair
8	128	Second in pair
9	256	Not primary alignment
10	512	Read fails platform/vendor quality checks
11	1024	Read is PCR or optical duplicate
12	2048	Supplementary alignment
Sum	0	

Decimal	Description of <b>second</b> read
1	Read paired
2	Read mapped in proper pair
4	Read unmapped
8	Mate unmapped
16	Read reverse strand
32	Mate reverse strand
64	First in pair
128	Second in pair
256	Not primary alignment
512	Read fails platform/vendor quality checks
1024	Read is PCR or optical duplicate
2048	Supplementary alignment
0	

### Common flags\*

One of the reads is unmapped:  
[73](#), [133](#), [89](#), [121](#), [165](#), [181](#), [101](#), [117](#),  
[153](#), [185](#), [69](#), [137](#)

Both reads are unmapped:  
[77](#), [141](#)


Mapped within the insert size and in  
correct orientation:  
[99](#), [147](#), [83](#), [163](#)

Mapped within the insert size but in  
wrong orientation:  
[67](#), [131](#), [115](#), [179](#)

Mapped uniquely, but with wrong insert  
size:  
[81](#), [161](#), [97](#), [145](#), [65](#), [129](#), [113](#), [177](#)

\* Collected from [here](#)

# Transcriptome Reconstruction



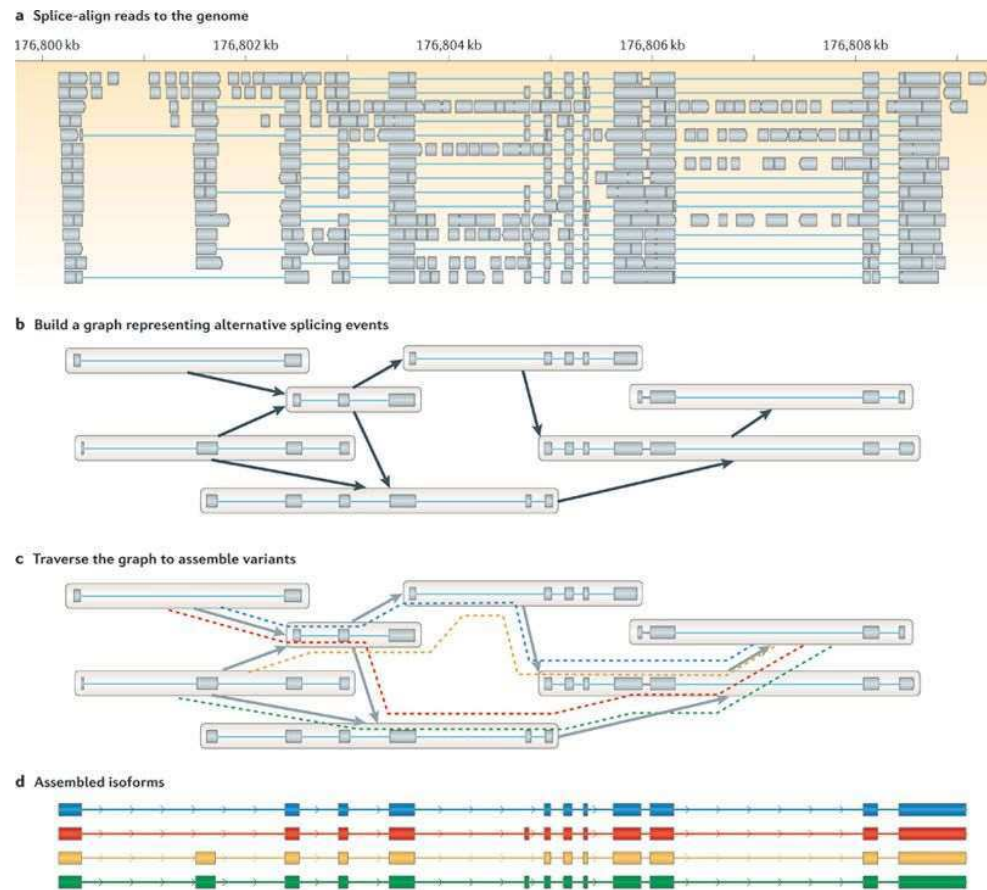
- This is an optional step in the process and I actually don't perform it very often
- You need paired-end reads and deep coverage
- Software:
  - Cufflinks
  - Stringtie
  - Scripture
  - Traph
  - Trinity
  - TransAbyss
  - Velvet

Reference guided assemblies

De novo assemblies

# Reference Guided Assemblies

- Splice junctions are the main guiding force in the reconstruction
- See which exons of these splice junctions overlap
- Estimate your isoforms



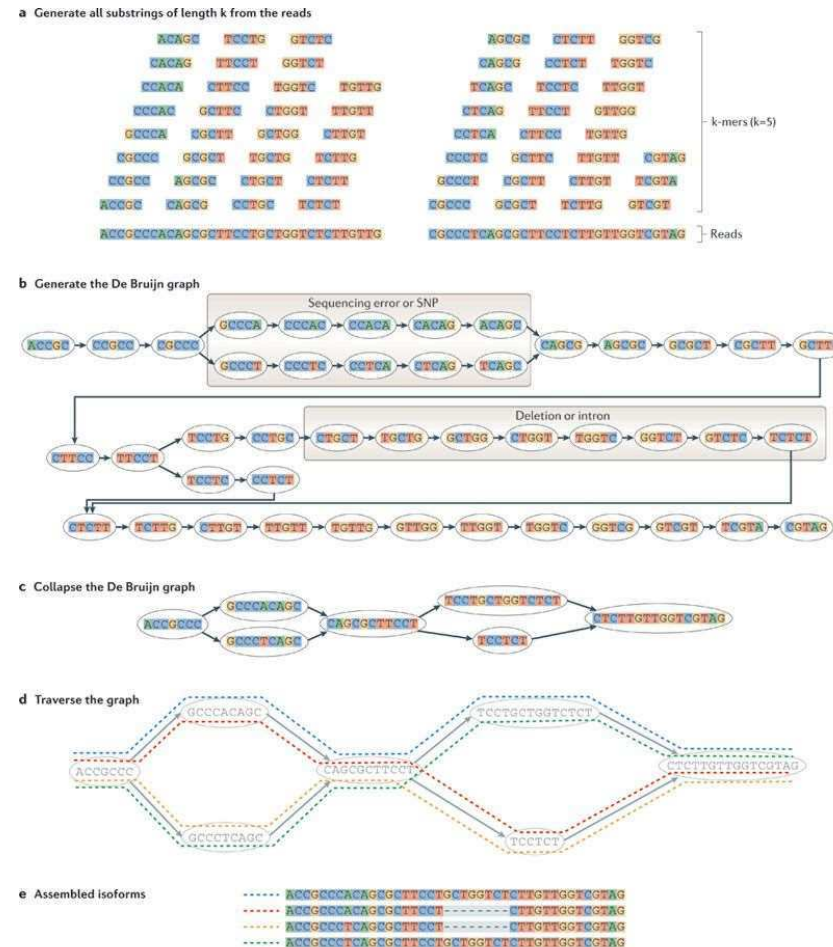
Nature Reviews | Genetics

Source: Martin & Wang 2015



# De Novo Assemblies

- Depends on De Bruijn graph
- Combinatorial mathematics, creates every possible combinations of length  $k$
- Identify indels and SNPs/errors
- Estimate your isoforms
- Computationally much more time than reference based approach



Nature Reviews | Genetics

Source: Martin & Wang 2015



# Gene Transfer Format (GTF)

- Transcriptome file format
- Can get references from ENSEMBL: Gene sets
- Easy to look at using R/rtracklayer

★	Species	DNA (FASTA)	cDNA (FASTA)	CDS (FASTA)	ncRNA (FASTA)	Protein sequence (FASTA)	Annotated sequence (EMBL)	Annotated sequence (GenBank)	Gene sets
Y	Human <i>Homo sapiens</i>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">EMBL</a>	<a href="#">GenBank</a>	<a href="#">GTF</a> <a href="#">GFF3</a>
Y	Mouse <i>Mus musculus</i>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">FASTA</a>	<a href="#">EMBL</a>	<a href="#">GenBank</a>	<a href="#">GTF</a> <a href="#">GFF3</a>

```
library(rtracklayer)
library(GenomicRanges)

isoref <- import("Y:/RNAseqProcessing/data/brain/HRDP.processing/getQualityTranscripts/gffcmp.annotated.cleaned.gtf")
isoref <- as.data.frame(isoref)

head(isoref)
```

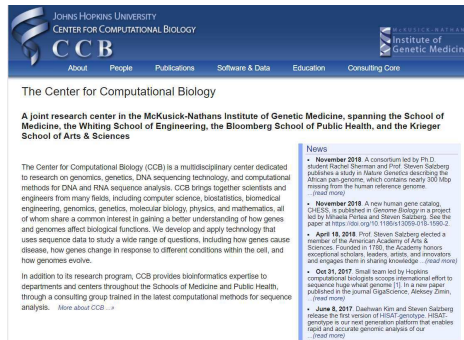
```
## seqnames start end width strand source type score phase
## 1 1 12358 13306 949 + StringTie transcript NA NA
## 2 1 12358 13306 949 + StringTie exon NA NA
## 3 1 12639 13301 663 + StringTie transcript NA NA
## 4 1 12639 12834 196 + StringTie exon NA NA
## 5 1 13255 13301 47 + StringTie exon NA NA
## 6 1 12914 13309 396 + StringTie transcript NA NA
## transcript_id gene_id xloc class_code tss_id exon_number
## 1 MSTRG.2.1 MSTRG.2 XLOC_000001 u TSS1 <NA>
## 2 MSTRG.2.1 MSTRG.2 <NA> <NA> <NA> 1
## 3 MSTRG.2.2 MSTRG.2 XLOC_000001 u TSS2 <NA>
## 4 MSTRG.2.2 MSTRG.2 <NA> <NA> <NA> 1
## 5 MSTRG.2.2 MSTRG.2 <NA> <NA> <NA> 2
## 6 MSTRG.2.3 MSTRG.2 XLOC_000001 u TSS3 <NA>
## gene_name ref_gene_id cmp_ref contained_in
## 1 <NA> <NA> <NA> <NA>
## 2 <NA> <NA> <NA> <NA>
## 3 <NA> <NA> <NA> <NA>
## 4 <NA> <NA> <NA> <NA>
## 5 <NA> <NA> <NA> <NA>
## 6 <NA> <NA> <NA> <NA>
```

```
isoref.oldschool <- read.delim(file="Y:/RNAseqProcessing/data/brain/HRDP.processing/getQualityTranscripts/gffcmp.annotated.cleaned.gtf", sep="\t", header=FALSE)
```

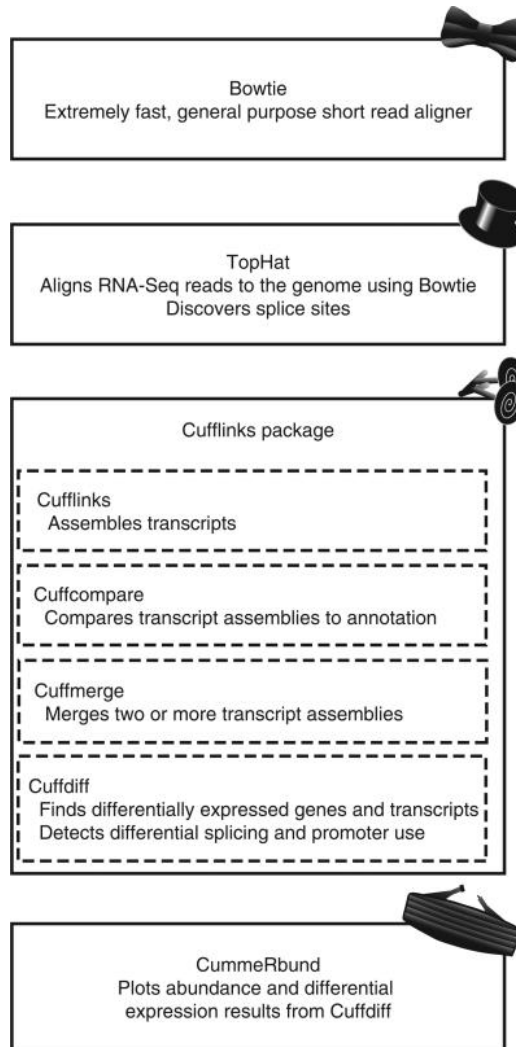
```
head(isoref.oldschool)
```

```
## V1 V2 V3 V4 V5 V6 V7 V8
## 1 1 StringTie transcript 12358 13306 . + .
## 2 1 StringTie exon 12358 13306 . + .
## 3 1 StringTie transcript 12639 13301 . + .
## 4 1 StringTie exon 12639 12834 . + .
## 5 1 StringTie exon 13255 13301 . + .
## 6 1 StringTie transcript 12914 13309 . + .
##
## v9
## 1 transcript_id MSTRG.2.1; gene_id MSTRG.2; xloc XLOC_000001; class_code u; tss_id TSS1;
## 2 transcript_id MSTRG.2.1; gene_id MSTRG.2; exon_number 1;
## 3 transcript_id MSTRG.2.2; gene_id MSTRG.2; xloc XLOC_000001; class_code u; tss_id TSS2;
## 4 transcript_id MSTRG.2.2; gene_id MSTRG.2; exon_number 1;
## 5 transcript_id MSTRG.2.2; gene_id MSTRG.2; exon_number 2;
## 6 transcript_id MSTRG.2.3; gene_id MSTRG.2; xloc XLOC_000001; class_code u; tss_id TSS3;
```

# Tuxedo Suite



This Johns Hopkins group produces all the “Tuxedo Tools”




New version of tophat is hisat2

Newer version of cufflinks is stringtie


This suite is extremely popular

Source: Trapnell et al, 2012

# Quantitate Reads



- Summarize on either the gene level or isoform level
- HTSeq (python module) or RSEM
- HTSeq not good for overlapping reads (i.e. not good for isoforms)



# RSEM Quantitation

- RNA-Seq Expectation Maximization (RSEM). Li et al, 2010
- Say you have a gene with 2 isoforms. Each has it's own unique exon

```
(isoform 1) AAAAAAAAAA  
(isoform 2) UUUUUUUUUU
```

- Say you have 3 reads for these

```
(read 1)  AAAAAAA  
(read 2)  UUUUUUU  
(read 3)  AAAAAAA
```

- Would say 2/3 of your reads for this gene come from isoform 1

# RSEM Quantitation 2

- However, example was extremely simplified as isoforms are highly similar and many don't have completely unique exons.
  - Isoform 1: Exon1-Exon2-Exon3-Exon4
  - Isoform 2: Exon1-Exon3-Exon4
  - Isoform 3: Exon2-Exon3-Exon4
- Notation
  - N = total # of reads (library size)
  - M = # known isoforms
  - L = read length
  - $l_i$  = length isoform
  - $\tau_i$  = TPM (fraction of transcripts belong to isoform i out of all transcripts in sample \* million)
  - $\theta_i$  = prior probability any single read derived from isoform i
- Take away  $\theta$  has a constant (uniform) probability density

$$\theta_i = \frac{\tau_i * l_i}{\sum_{k=1}^M \tau_k * l_i}$$

$$\tau_i = \frac{\theta_i / l_i}{\sum_{k=1}^M \theta_k / l_i}$$


$$\theta_i \propto \tau_i * l_i$$

$$\sum_i \theta_i = \sum_i \tau_i = 1$$

$$p(\theta) \propto 1$$

# RSEM Quantitation 3

- What does an estimate of  $\Theta$  (or  $\tau$ ) look like?
- Full posterior distribution is  $p(\Theta | R)$  or  $p(\tau | R)$
- Li et al 2010, describe following Bayesian network (a probabilistic graphical model):



- $G_n$  is isoform and  $S_n$  is starting position
- Probability of a read coming from a specific gene with influence the number of reads you see in your data

# RSEM Quantitation 4

Start with a joint probability:

$$p(G, S, R|\theta) = \prod_{n=1}^N p(G_n, S_n, R_n|\theta) = \prod_{n=1}^N p(G_n|\theta)p(S_n|G_n)p(R_n|G_n, S_n).$$

With lots of math:

$$p(\theta|R) = \frac{p(R|\theta)p(\theta)}{p(R)}.$$

Start an estimate of  $\theta_1$  and  $\theta_2$  (assuming a simple 2 isoform dataset with  $N_1$  # reads aligning uniquely to isoform 1 and  $N_2$  # reads aligning uniquely to isoform 2

$$\begin{aligned}\theta_1^{(0)} &= N_1 / (N_1 + N_2), \\ \theta_2^{(0)} &= N_2 / (N_1 + N_2).\end{aligned}$$

# RSEM Quantitation 5

- Start doing expectation maximization for finding the maximum a posteriori estimate of  $\Theta$

- $N_{12}$  = # reads overlapping both isoforms
  - $N$  = # read total ( $N_1 + N_2 + N_{12}$ )

$$\theta_i^{(1)} = \frac{N_i + N_{12} \cdot \tau_i^{(0)}}{N}$$

- Repeat this cycle till  $\theta^{(r+1)}$  does differ too much from  $\theta^{(r)}$
- Simulation example:

Truth

$i$	$l_i$	$\theta_i$
1	300	0.60
2	1000	0.10
3	2000	0.30

Counts

$N_1 = 111$	$N_{12} = 69$	$N_{123} = 144$
$N_2 = 26$	$N_{13} = 311$	
$N_3 = 186$	$N_{23} = 153$	

RSEM iterations

$i$	$\theta_i^{(0)}$	$\theta_i^{(1)}$	$\theta_i^{(2)}$	$\theta_i^{(3)}$	$\theta_i^{(4)}$
1	0.34	0.53	0.58	0.59	0.59
2	0.08	0.07	0.07	0.08	0.08
3	0.58	0.40	0.34	0.33	0.32



# RSEM Code

- Like alignment tools, you need to build a reference
- It is actually aligning to the **transcriptome** within this process (bowtie)


```
rsem-prepare-reference --gtf referenceTranscriptome.gtf --bowtie2 referenceGenome.fa /pathToIndexOutput/suffix
```

- RSEM on sample:

```
rsem-calculate-expression -p 8 --time --seed 2020 --bowtie2 --paired-end --seed-length 20 /data/hi-seq/MuKO.Brain.Mouse/alignedReads/HISAT2/rRNA/MuKOHet_2.rRNA/sample1.end1.fq /data/hi-seq/MuKO.Brain.Mouse/alignedReads/HISAT2/rRNA/MuKOHet_2.rRNA/sample1.end2.fq /data/hi-seq/MuKO.Brain.Mouse/quantitation/MuKOHet/RSEM.ensembl MuKOHet_2
```

# RSEM Batch

- Python script called: runRSEM\_batch.py
- Written by Spencer Mahaffey (thank you again!)

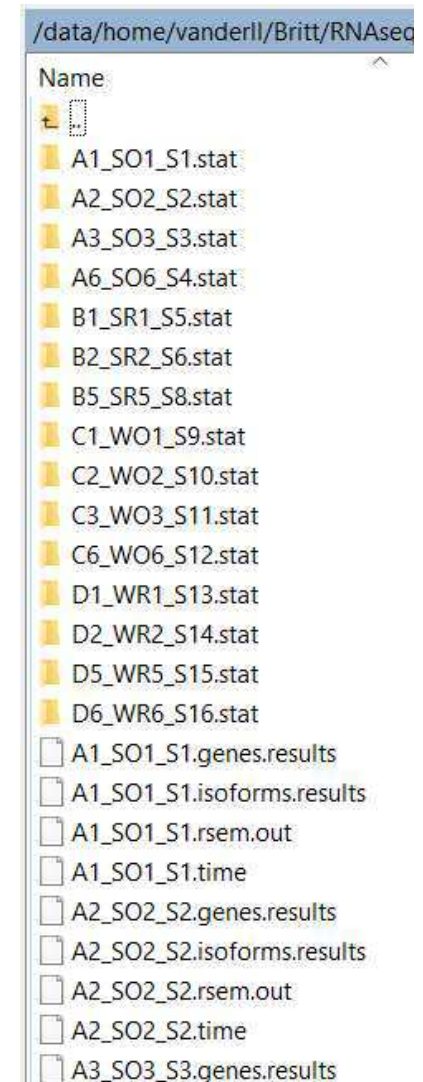


Spencer Mahaffey  
smahaffey

```
python runRSEM_batch.py --rsem-time --rsem-seedLen 20 --rsem-seed 2020 --rsem-bowtie2 --rsem-noBam -  
-rsem-fwProb 0.0 --paired -d _R -o /data/home/vanderll/Britt/RNAseq.20171106/RSEMquant.ENSEMBL/ -i tr  
immed.fastq.gz /data/home/sabal/Britt/RNA-Seq.2017-11-06/trimmedReads/ /data/home/vanderll/Britt/RNAs  
eq.20171106/index/ dm6.ensembl 8
```

# RSEM Output

- Folder with sample name:
  - 3 files regarding mathematical files
  - 1 files for count estimated at each iteration
  - 1 file for other mathematical estimates in model
- 4 Files per sample:
  - sampleName.genes.results
  - sampleName.isoforms.results
  - sampleName.rsem.out
  - sampleName.time



# RSEM Output 2

- The genes.results and isoforms.results files have the same format

```
rsem <- read.table(file="Y:/Britt/RNAseq.20171106/RSEMquant.ENSEMBL/A1_S01_S1.isoforms.results", sep=
"\t", header=TRUE)
head(rsem)
```

##	transcript_id	gene_id	length	effective_length	expected_count	TPM
## 1	FBtr0081624	FBgn0000003	299	99.47	981.00	772.69
## 2	FBtr0071763	FBgn0000008	4847	4644.16	309.96	5.23
## 3	FBtr0071764	FBgn0000008	5173	4970.16	19.51	0.31
## 4	FBtr0100521	FBgn0000008	4665	4462.16	200.53	3.52
## 5	FBtr0342981	FBgn0000008	3897	3694.16	0.00	0.00
## 6	FBtr0083387	FBgn0000014	4458	4255.16	0.00	0.00
##	FPKM	IsoPct				
## 1	656.75	100.00				
## 2	4.44	57.73				
## 3	0.26	3.40				
## 4	2.99	38.87				
## 5	0.00	0.00				
## 6	0.00	0.00				

# RSEM Merge Data

```
wd = "Y:/Britt/RNAseq.20171106/RSEMquant.ENSEMBL/"

getGeneResults = function(a){
  b = a[grep(".gene",a)]
  return(b)
}

files = paste(wd, getGeneResults(list.files(wd)), sep="")

#load in the data
for(i in 1:nrow(files.v2)){
  x = read.table(file=files.v2[i,"file"],sep="\t",header=TRUE)
  x = x[,c("gene_id","expected_count")]
  colnames(x)[2] = files.v2[i, "sample"]
  if(files.v2[i,"file"]!=files.v2[1, "file"]) rsem = merge(x,rsem,by=c("gene_id"),all=TRUE)
  if(files.v2[i, "file"]==files.v2[1, "file"]) rsem = x
}

x = read.table(file=files.v2[i,"file"],sep="\t",header=TRUE)

estCnts = rsem[,-1]
rownames(estCnts) = rsem$gene_id

counts = round(estCnts,0)
save(counts, file="Y:/Britt/RNAseq.20171106/data/RSEM.ensembl.estCounts.noFiltering.Rdata")
```

# Reporting Preprocessing

- Data
  - When downloaded
  - Where saved
- Processing code
  - Just referencing where or what was run
  - Versions of programs
- Summary of reads at each step
- If quantitating to a known reference transcriptome, like to check out percentage Ensembl “biotype” aligned to
- Look at Rmarkdown report
- The Rmd and html will be available on yampa under:

# References

Hansen KD, Brenner SE, Dudoit S (2010). *Biases in Illumina transcriptome sequencing caused by random hexamer priming*. Nucleic Acids Res. 38(12):e131.

Pertea M, Pertea GM, Antonescu CM, Chang TC, Mendell JT, Salzberg SL. (2015) *StringTie enables improved reconstruction of a transcriptome from RNA-seq reads*. Nat Biotechnol. 33(3):290-5.

Martin JA, Wang Z. (2011) *Next-generation transcriptome assembly*. Nat Rev Genet. 12(10):671-82.

Cole Trapnell, Adam Roberts, Loyal Goff, Geo Pertea, Daehwan Kim, David R Kelley, Harold Pimentel, Steven L Salzberg, John L Rinn, and Lior Pachter (2012) *Differential gene and transcript expression analysis of RNA-seq experiments with TopHat and Cufflinks*. Nat Protoc. 2012 Mar 1; 7(3): 562–578.

Li, B., Ruotti, V., Stewart, R.M., Thomson, J.A. & Dewey, C.N. (2010). *RNA-Seq gene expression estimation with read mapping uncertainty*. Bioinformatics **26**, 493–500.

<https://www.biostat.wisc.edu/bmi776/lectures/rnaseq.pdf>